PEEL - User Manual

Pierre Terdiman – v1.1
pierre.terdiman@gmail.com

1) Introduction

Welcome to the Physics Engine Evaluation Lab - a.k.a. PEEL.

PEEL is a tool designed to evaluate, compare and benchmark physics engines. In a way, it is very similar to the old PAL project (Physics Abstraction Layer).

It was initially written to compare PhysX versions between each-other, and catch performance regressions. Support for entirely different engines was added later, giving the tool a much larger scope. To this date it has been successfully used by various people to discover previously unknown issues in their engines, and actually improve them.

2) Physics engines selection

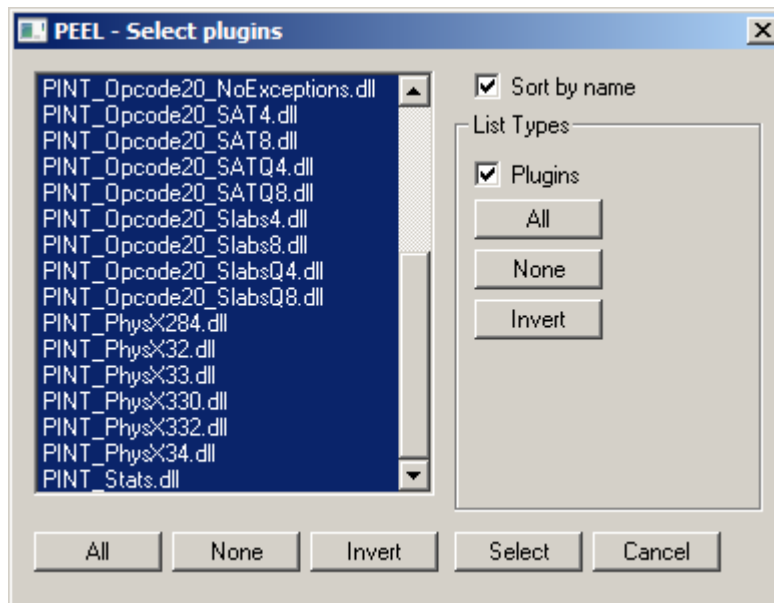When you run PEEL, the first thing you see is the plugin selection dialog:

Figure 1: plugin selection dialog

Each plugin is a DLL whose name starts with "PINT" (for Physics INTerface). There is one plugin for each physics engine. To select a plugin, click on it so that it is highlighted in blue, then click on the 'Select' button to validate your selection. You can choose as many plugins as you want here. Otherwise if you are only interested in a single plugin, you can simply double-click on its name directly.

Note: there is an internal limit of 32 plugins, and the Fx function keys are used to enable/disable an engine. So it is best to select less than 12 engines at once.

3) **Main window – Menu Mode**

All selected plugins are loaded. If the loading is successful, PEEL's main window should appear next. This is called the *Menu Mode*:
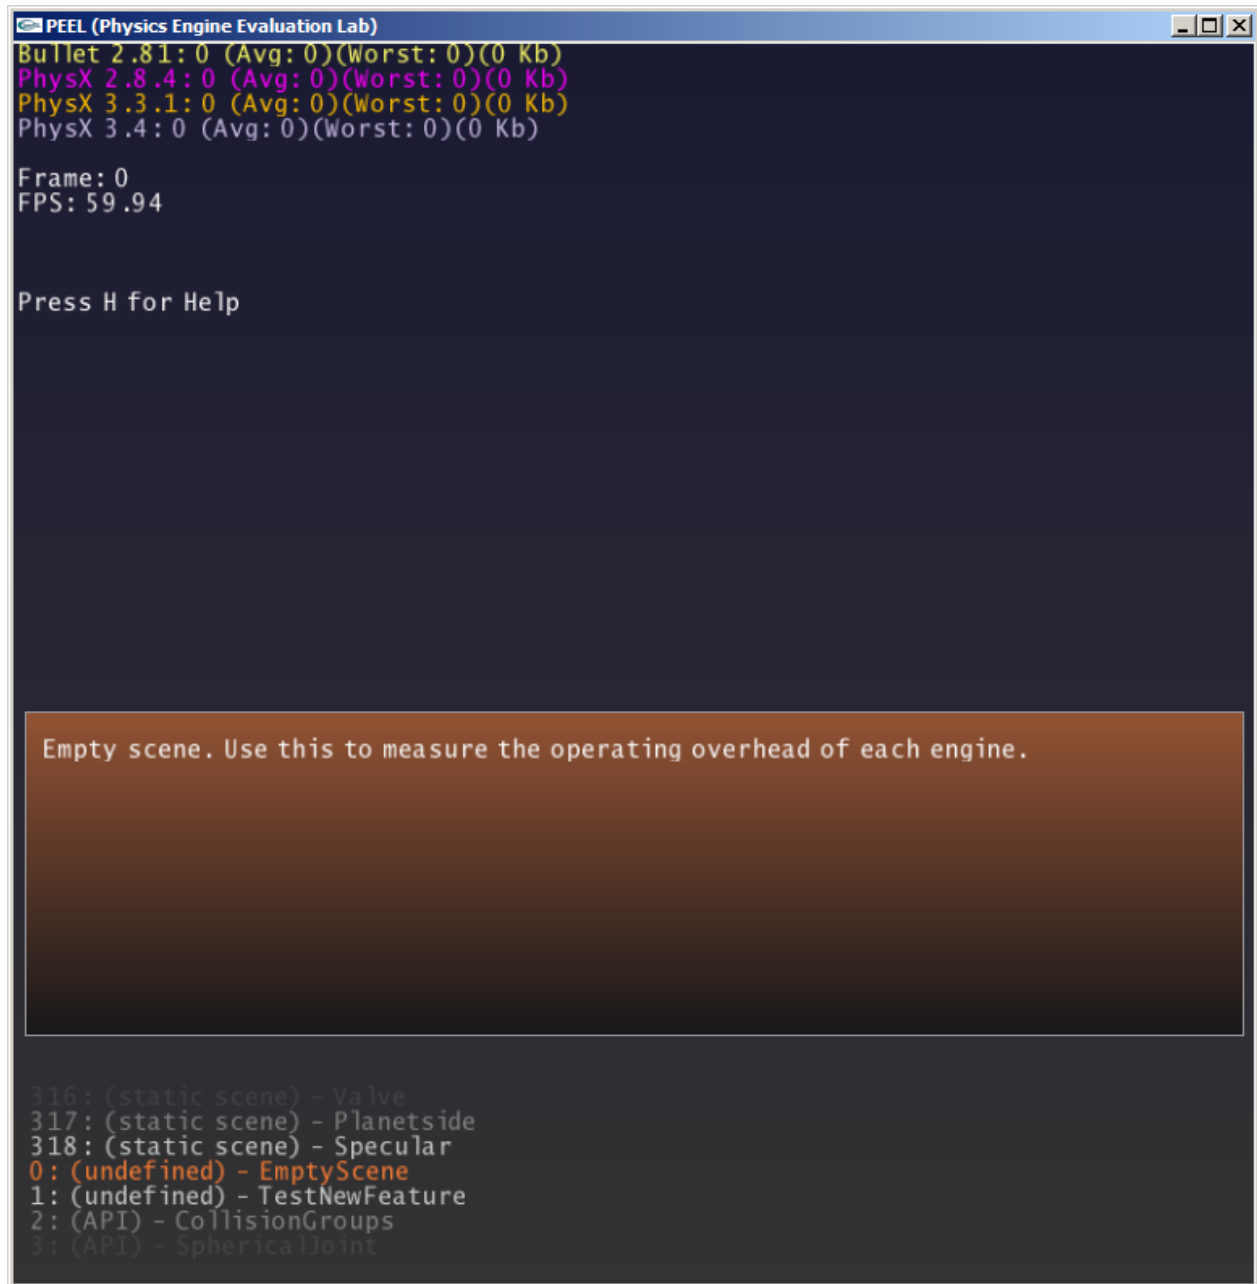

Figure 2: Main window, Menu Mode

The names of loaded plugins are listed at the top of the screen. In Figure 2 for example, 4 different plugins have been loaded. The numbers following the plugins' names are performance and memory stats for each of them.

Test selection

At the very bottom of the screen there is a list of available tests. Each test has the following format:

Index: (category) – name

Index is simply the test number. It has no purpose other than knowing the total number of available tests.

Category is the group to which the test belongs.

Name is the test's name.

Each test also has a description string, displayed in the brown window located above the test selection area.

Use the arrow keys to navigate & select desired test. The Up and Down arrow keys go to the previous or next test in the list. The Left and Right arrow keys skip an entire category, which makes the navigation quicker when you know what category you are interested in.

Once the desired test is selected (highlited in orange), press the Return key to activate the test and switch the window to *Simulation Mode.*

Categories

Tests are roughly sorted into the following categories:

- Undefined: tests that do not fall into one of the other categories.

- API: tests that check the API works as expected. Usually very simple scenes doing only one specific thing (e.g. checking that collision filtering works)

- Behavior: tests that check the physics behavior is as it should be, or at least plausible. Can be seen as correctness tests, as opposed to performance tests.

- Contact generation: this category is dedicated to contact generation problems and pitfalls.

- Joints: this category contains scenes that use traditional joints like spherical, hinge, prismatic, etc. These joints typically use the same iterative solver as the contacts.

- Articulations: some physics engines support more advanced or more stable joints dedicated to articulated systems. They are usually more expensive than the regular joints, and may use dedicated solvers. This category is for them.

- Performance: generic performance tests usually involving a lot of rigid bodies or complicated scenes.

- Kinematics: tests that use kinematic objects.

- Vehicles: tests that use dedicated, vehicle-specific modules.

- Raycast: mainly performance tests for raycasts. Might include correctness tests as well.

- Sweep: mainly performance tests for sweeps. Might include correctness tests as well.

- Overlap: mainly performance tests for overlaps. Might include correctness tests as well.

- CCD: tests dedicated to continuous collision detection problems.

- Static scene: tests that only load a static scene, no rigid bodies. Mainly used for raytracing benchmarks.

- WIP: work-in-progress tests.

Note that these categories are actually overlapping quite a bit, and thus the sorting is only approximate. For example several tests in the kinematics category are only concerned about the performance of kinematics objects, and thus they would also fit in the performance category.

4) <u>Main window – Simulation Mode</u>

In simulation mode, the selected test runs for all engines at the same time, and the results are rendered on the same screen. Each engine has a main color used both to display its name at the top of the screen, and the rigid bodies it controls. Rendering all engines to the same screen allows one to easily spot differences in behaviours between them. From a performance point of view, it also emulates to a small extent how the physics engine would be used in a real game, where other subsystems would run before and after the physics simulation is performed. Benchmarking the engines that way might provide more realistic results than benchmarking them in isolation of everything else.

The top line contains the name of current test, and the number of available camera positions. Most tests only have one camera position available, but some of them (especially raytracing tests) have many. Use the + and – keys to change the camera position.

For each physics engine there are 4 performance numbers updated each frame. The first number, just after the physics engine's name, is the simulation time for current frame. *(Avg: x)* is the average simulation time. *(Worst: x)* is the largest simulation time so far. The last number is the memory used by each engine, provided the allocations have all been properly re-routed to the engine's user-defined allocators.

Frame: x is a simple frame counter. It starts counting when entering Simulation Mode, and stops counting when returning to Menu Mode. It is possible to pause/resume the simulation by pressing P, and do frame-by-frame simulation by pressing O (for "One frame").

FPS: x is a simple framerate counter. This is not terribly interesting in itself, since it includes the accumulated simulation time of each engine, as well as the rendering time. It is possible to disable rendering entirely by pressing R. It is also possible to enable/disable each physics engine by pressing the appropriate function keys. F1 maps to the first listed engine, F2 to the second, etc. The framerate counter might be capped to 60 Hz if VSync is enabled. It is possible to enable/disable VSync in PEEL's Options Panel.

While a test is running, press Return to stop the simulation and return to Menu Mode.

Press S to save an Excel file containing the current test's performance data so far. PEEL currently does not display performance graphs on its own, so it relies on Excel to create these.

Move the camera with the arrow keys and the mouse (using left-button drag). Reset the camera position by pressing C. The camera position is not automatically reset when running the same test twice. This allows one to move the camera to an interesting spot, stop the simulation by going to Menu Mode, and start the simulation again from the desired point of view by starting the same test again. An alternative is

to pause the simulation in Menu Mode, start the simulation paused, move the camera to the desired spot, then unpause.

In simulation mode, use the right mouse button to do an action. Select desired action in the Options panel's generic tab, under "Current tool". Default action is picking/manipulating rigid bodies. It is possible to delete the currently picked/manipulated object by pressing the "Suppr" key while the object is captured by the mouse.


Figure 3: Main window, Simulation Mode

5) **Per-test UI (configurable tests):**

Since PEEL 1.1 each test can have its own UI panel, which appears when the test is selected and simulation mode starts. These configurable tests often gather together several of the PEEL 1.0 tests, which explains why the total number of tests reported in Menu Mode is smaller in PEEL 1.1 (even though it actually contains more tests overall than the previous version).

Sometimes the settings used in the PEEL 1.0 scenes have been captured in "preset" scenes, like for example in the ragdolls test scene:

Using the presets will recreate the same scenes as in PEEL 1.0.

Otherwise these dialogs are pretty much self-explanatory: tweak the parameters as you wish, and press the "Reset test" button to restart the test with the new settings. Alternatively, go back to the Menu Mode and activate the same test again: the new settings will then be taken into account. The settings are not reset until you start a different test.

6) **Options panel:**

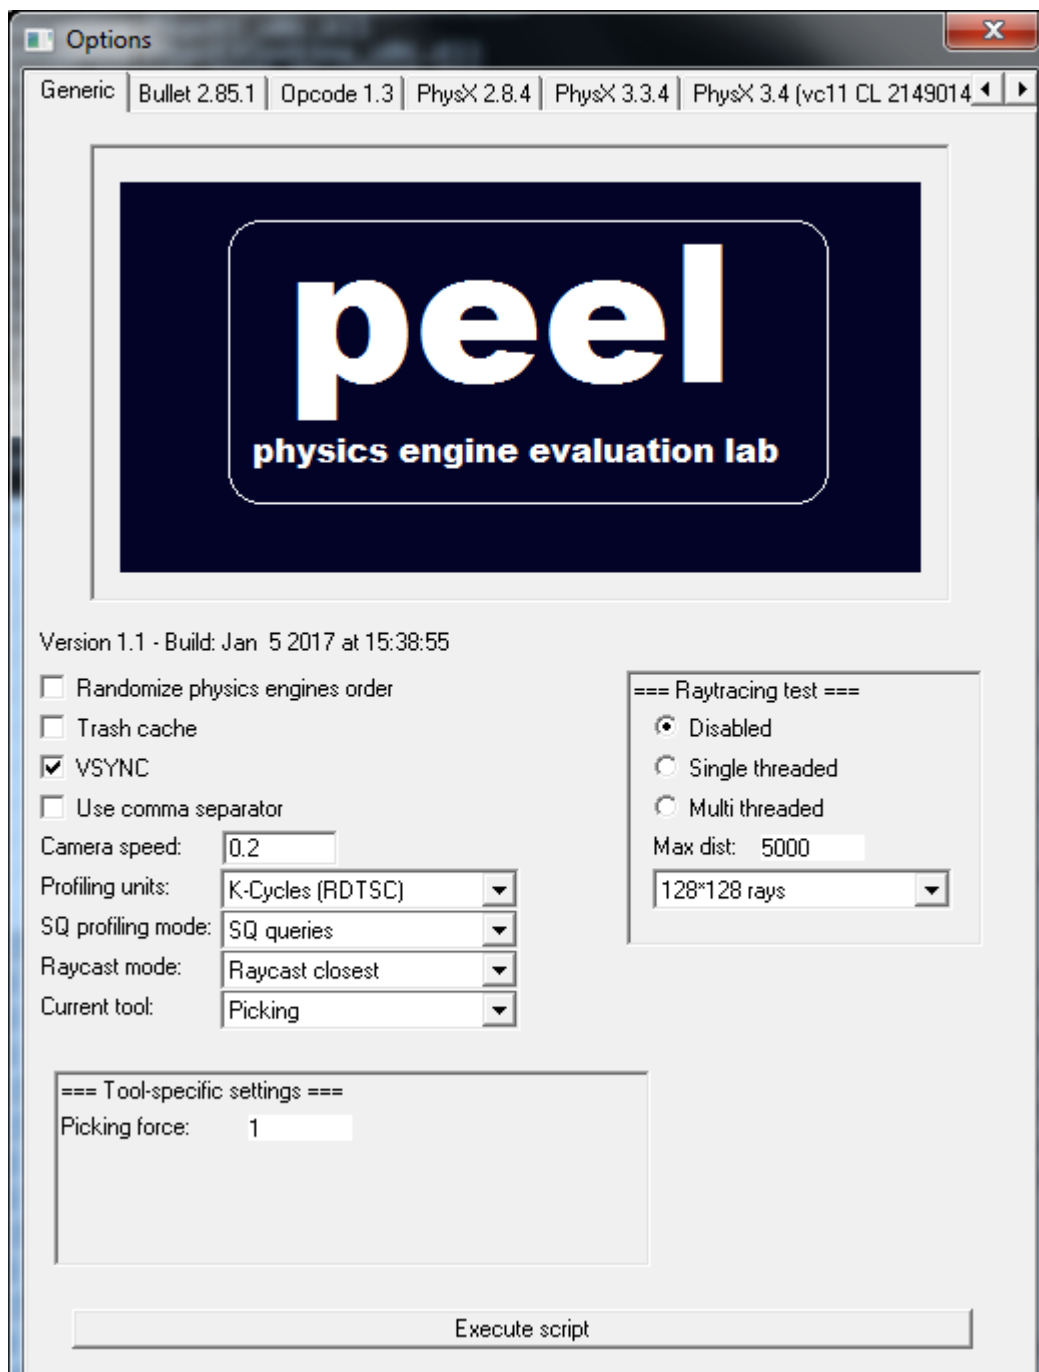The options panel has one tab for shared/generic options, and one tab for each loaded plugin.

Figure 4: Options Panel – shared (generic) options

The generic tab contains settings affecting all engines:

Randomize physics engines order

In Simulation Mode, all engines run sequentially. By default they run in the order in which they are listed at the top of the screen, which depends on the plugins' DLL names. There was a concern that the order may have an impact on performance results, so this feature randomizes the order each frame to make sure that benchmarks are not biased.

Trash cache

The state of the cache has a large impact on performance. Running a small scene in a single engine, in isolation from the actual game engine, possibly without rendering, is a best case scenario that is likely to produce unrealistic performance figures. The same scene in the actual game will probably have lower performance. Running PEEL with multiple physics engines at once and rendering enabled already takes care of this issue to some extent. But this feature goes a bit further and actually tries to flush the cache each frame, to emulate what might happen in an actual game. This is rather experimental so far.

VSYNC

This feature simply enables or disables VSYNC.

Use comma separator

Defines which character is used as a separator in saved Excel files (',' or ';').

Camera speed

In Simulation Mode it is possible to move the camera with the arrow keys. This parameter defines the camera's speed when doing so.

Profiling units

This feature defines the units of reported performance numbers. Available options are:
- K-Cycles (RDTSC): 1 reported unit = 1024 cycles, as measured by the RDTSC instruction.
- ms (timeGetTime): 1 reported unit = 1 ms, as measured by the timeGetTime function.
- us (QPC): 1 reported unit = 1 us, as measured by the QueryPerformanceCounter function.

SQ profiling mode

SQ stands for Scene Queries, and an "SQ test" is one that focuses on raycast, sweep or overlap queries. Physics engines often have two very different codepaths here: one for managing the SQ structures (creating them or updating them when rigid bodies move each frame), and one for actually performing raycast, sweep or overlap queries against the structure. Both of these codepaths are important and worth profiling. This feature defines which part of the physics engines should be profiled:
- Simulation: profiles the structure's management (creation/update/etc), which usually happens during the rigid body simulation.
- SQ queries: profiles the actual raycast, sweep and overlap queries.
- Combined: profiles both and combines the results.

Note that this is also used for tests that add objects to the scene at runtime. The "SQ queries" mode in that case would profile the cost of the "add" calls. But this is only half of the story: sometimes the actual work only happens during the next simulate call, and its cost would be captured in the "Simulation" or "Combined" modes.

Raycast mode

This feature defines what kind of raycast queries are used in the "Raycast" test category:

- Raycast closest: these are the regular raycasts returning a single hit (the closest one along the ray).

- Raycast any: these are boolean raycasts simply returning whether the ray touched something or not. They should be faster than "raycast closest" queries.

- Raycast all: these are raycasts returning all objects touched by the ray. They should be slower than "raycast closest" queries.

It would have been possible to create dedicated raycast tests for each of these modes, but it would have potentially multiplied the total number of raycast tests by 3 (to cover all cases). Instead, selecting the mode as a global option reduces the total number of tests. On the other hand it is slightly more tedious to use.

Current tool & tool-specific settings

This feature defines what happens when pressing the right mouse button (RMB) in Simulation Mode. Available tools are:

- Picking: users can select and drag rigid bodies using the RMB.
- Add impulse: pressing the RMB applies an impulse at the picked point on the picked object.
- Shoot box: a new box is created, added to the scene, and shot in the view direction.
- Camera tracking: selects the object that the camera will be tracking. After selecting an object here, you can for example go back to the picking tool and move that object around with the mouse. The camera will follow it.
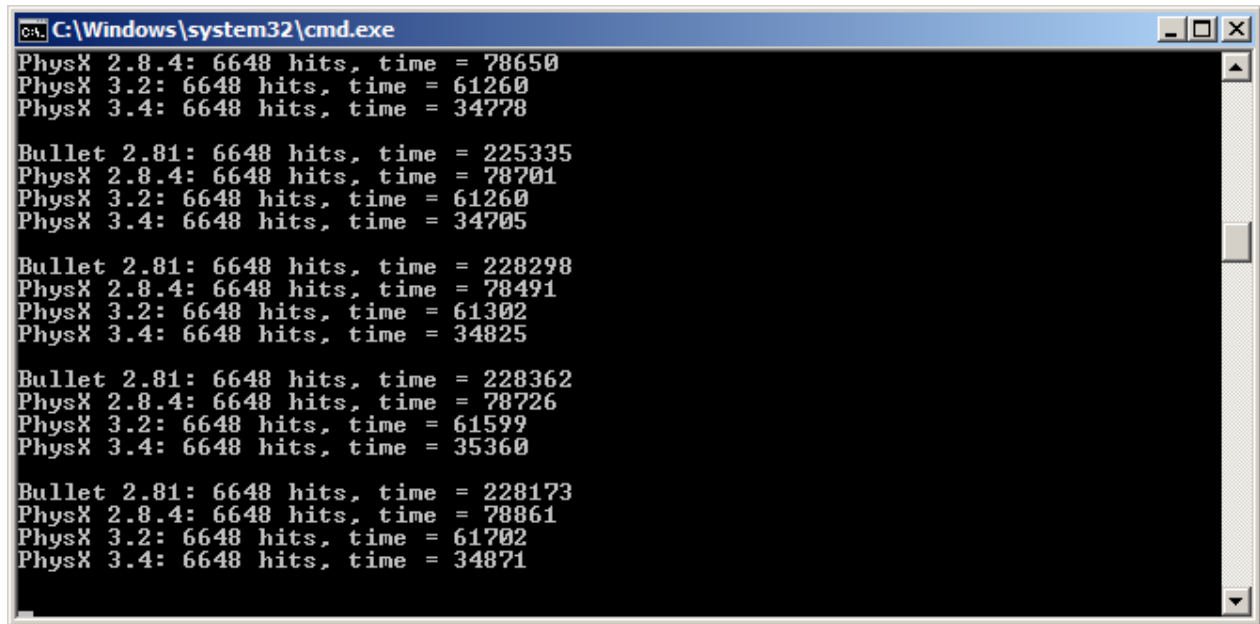
Each tool can have additional settings in the "tool-specific settings" area of the UI. For example the picking force for the picking tool can be customized there. The contents of this area will change according to selected tool.

Raytracing test

This feature uses the physics engines' raycast queries to raytrace the scene. The results are rendered in each plugin's dedicated tab, in the "Raytraced view" window. Shading is very basic, since the goal is not to create pretty pictures but to quickly visually check that everything is working as expected. The raytracing test has a few options of its own:

- Disabled: disables the raytracing test.
- Single threaded: runs the test with a single thread.
- Multi threaded: runs the test with 4 threads (each thread is given a quarter of the rendering window).
- Max dist: ray length.
- Number of rays: size of rendering window.

When the raytracing test is enabled, performance numbers are captured for the whole operation and printed in the DOS console. Warning: this is not the same as for the regular tests in PEEL, whose results are printed on screen. For raytracing tests the units are currently K-cycles and cannot be changed. The number of raycast hits is also reported to the DOS console (Figure 5).

```
C:\Windows\system32\cmd.exe                                                    _ □ ×
PhysX 2.8.4: 6648 hits, time = 78650
PhysX 3.2: 6648 hits, time = 61260
PhysX 3.4: 6648 hits, time = 34778

Bullet 2.81: 6648 hits, time = 225335
PhysX 2.8.4: 6648 hits, time = 78701
PhysX 3.2: 6648 hits, time = 61260
PhysX 3.4: 6648 hits, time = 34705

Bullet 2.81: 6648 hits, time = 228298
PhysX 2.8.4: 6648 hits, time = 78491
PhysX 3.2: 6648 hits, time = 61302
PhysX 3.4: 6648 hits, time = 34825

Bullet 2.81: 6648 hits, time = 228362
PhysX 2.8.4: 6648 hits, time = 78726
PhysX 3.2: 6648 hits, time = 61599
PhysX 3.4: 6648 hits, time = 35360

Bullet 2.81: 6648 hits, time = 228173
PhysX 2.8.4: 6648 hits, time = 78861
PhysX 3.2: 6648 hits, time = 61702
PhysX 3.4: 6648 hits, time = 34871
```

Figure 5: output of raytracing test

Execute script

This feature allows one to run a set of tests automatically via a script. A script is a simple text file with the following format and supported commands:

```
// Script name

NbFrames            100    // Default number of frames to simulate
Rendering           false  // Enable or disable rendering
RandomizeOrder      true   // Randomize engine order each frame, or not
TrashCache          false  // Trash cache after each simulation call, or not

// Test <name> <#frames>

Test SceneRaycastVsStaticMeshes_MeshSurface 1400
Test SceneRaycastVsStaticMeshes_TestZone_ShortRays
…
```

The script name and initial parameters should be obvious. The script then runs each test sequentially for the given number of frames. An Excel file containing performance data is automatically saved at the end of each test. The number of frames to simulate can be customized for each test, by writing the desired number after the test name (e.g. "1400" in the above script). If that number is not specified, the default number is used (e.g. "100" in the above script).

NOTE: there is currently an issue with configurable tests and scripts. It is currently not possible to configure the tests from the script, so the scripted tests will run using each test's default parameters. This means that it is not possible anymore to run some of the PEEL 1.0 preset scenes from a script.

Per-engine options

Each physics engine can have its own options, captured in a separate tab for each plugin/engine. These options are engine-specific and thus beyond the scope of this document, but they typically map to something exposed in the engine's API. Thus, please refer to each engine's documentation for details.

Note that some options are activated immediately in the running test, while others are only activated after the test is restarted. In doubt, always restart the test after changing the options in the UI.
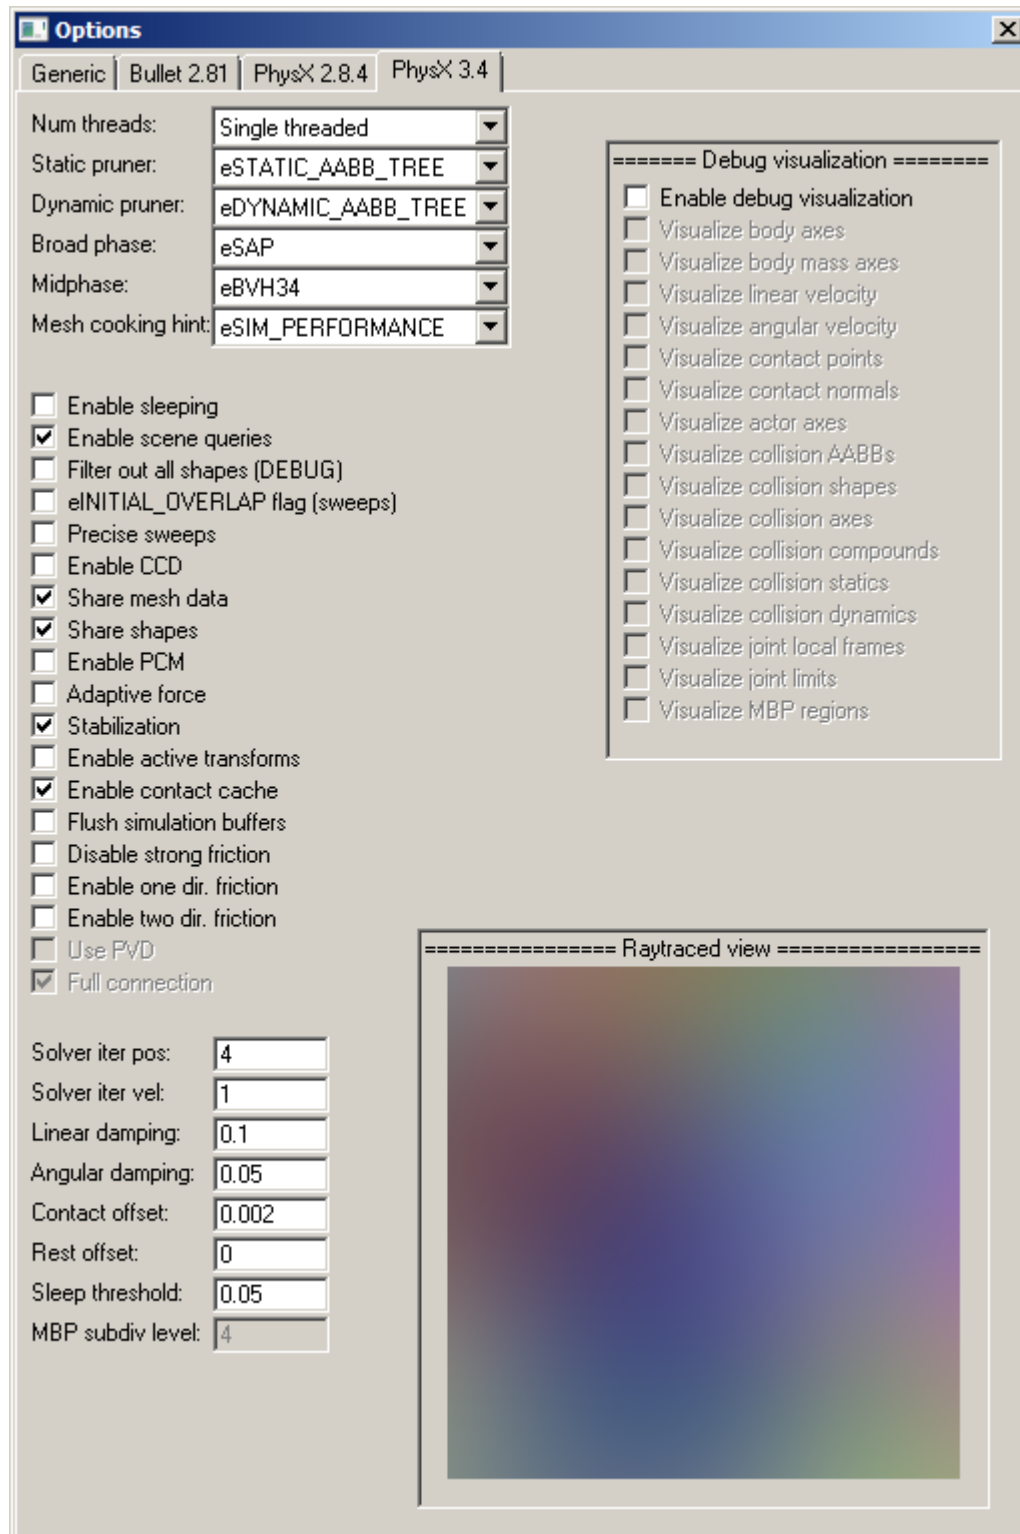

Figure 6: Options Panel – per-engine options

Scene stats

There is a special plugin named PINT_Stats.dll. This is not a physics engine. This is a module gathering scene statistics, like the number of created objects in a scene, and their type. PEEL itself does not gather and display such stats, but this plugin does.
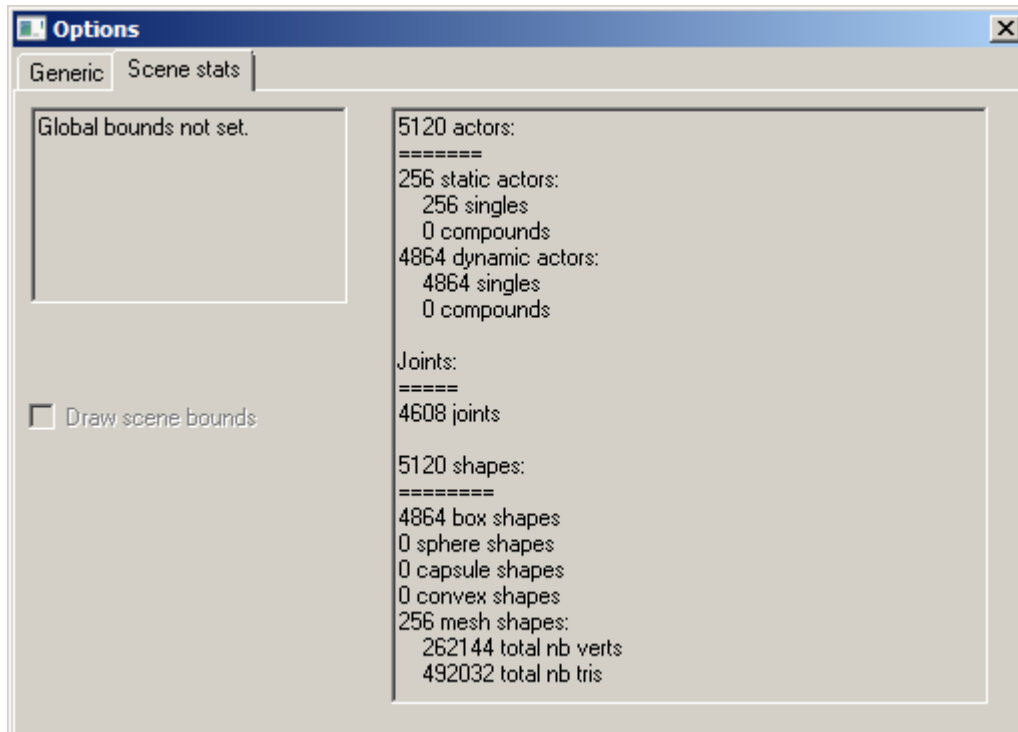
Figure 7:

7) Command line options

PEEL supports the following command line options:

-r	Disable rendering
-p filename.dll	Load given plug-in
-t test_name	Select given test
//	Disable all further command line arguments

This last one is handy if you want to temporarily ignore command line arguments but still keep them around in Visual Studio's property page. You can then have something like the following command arguments:

-r // -p PINT_Opcode20.dll -t SceneRaycastVsStaticMeshes_TessBunny16384

...and only the first one will be taken into account.

8) Typical PEEL usages

Regression testing

PEEL can be used to catch performance or behavioral regressions from one version to another.

For example when running several PhysX versions at the same time, runtime performance for a new version should always be better or at least equal to the one from the previous version. If it is not, then either there is a known change that explains the regression (and the regression might be acceptable), or it is a bug / unexpected side effect that must be fixed.

Similarly, if for example the friction appears to be very different in one version while the friction model did not change, it is probably a bug that was not caught by unit tests or visual inspection.

There are a lot of subtle things that do not appear to be immediately wrong when constantly working on the last version, but running the old and new versions at the same time on the same screen makes a lot of these things suddenly very obvious.


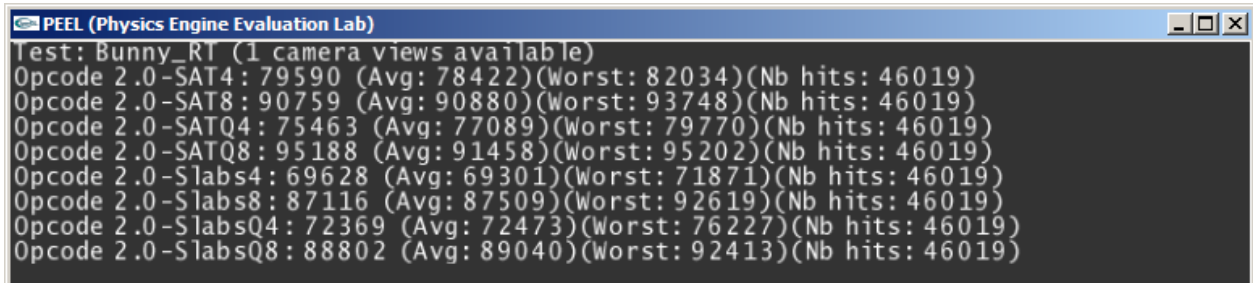Figure 8: PhysX performance going in the right direction in this scene

Figure 9: Excel graph from PEEL's recorded data. This makes things even clearer.

Benchmarking different options during development

PEEL can be used to benchmark different approaches while developping a new feature.

For example you might have a define in your code that enables one implementation or another, say e.g. SIMD vs non-SIMD. With PEEL it is easy to compile a PINT plugin using one implementation, change the PINT DLL's name, recompile with the other implementation, and run the two competing versions against each other. Since they run at the same time, they suffer from the same background noise at the same time, and the resulting performance curves are usually quite clear with one version always above the other. The winner is usually clearer with PEEL than with more traditional benchmarks running competing implementations sequentially.

Figure 10: different codepaths in Opcode 2.0. Which one is best?

Raycast debugging

PEEL can be used to debug raytracing functions.

It is very easy to break a raytracing function without noticing. PEEL's raytracing test can make subtle bugs visually obvious (unexpected black pixels appearing here and there in the raytracing window, backface culling or ray length not being properly respected, gaps appearing between objects that should touch, etc). It is then possible to trigger a breakpoint when e.g. a black pixel happens, and then just trace the code.


Figure 11: unit tests pass, yet visually it's clearly broken.

Benchmarking against the competition

PEEL can be used to benchmark engine A against engine B.

This is the purest way to do it, as comparing timings of two apps using different physics engines but also different architectures, renderers, or even different profiling functions is obviously dubious. PEEL uses the same tests with the same input values going through the same interface, with the same objects rendered the same way, profiled the same way, on the same machine, with the same background noise at the same time. This is the most objective way to compare engines. That being said, there are a lot of pitfalls here, explained in Appendix A.


Figure 12: different engines running the same test at the same time.

Note the spike in figure 12. It happened because an unrelated program started an expensive operation in the background. With PEEL, the unexpected event affects all 3 engines, and thus the recorded performance data remains fair and reliable. If we would run 3 tests sequentially, each one with a different physics engine, the spike would only happen during one of the tests, giving an unfair advantage to the other two. Incidentally, the amplitude of the spike reveals that some engines deal with external disturbances better than others.

Write once, run everywhere

PEEL can be used as a complement to the PhysX samples and migration guide from one version to the next. Or even from one engine to the other.

If you missed the simple samples from PhysX 2.8, PEEL's simple tests can be used as a replacement for the missing equivalent samples in PhysX 3.x. If the migration guide from PhysX 2.8 to PhysX 3.x is incomplete, PEEL might be able to fill the gap if it uses the feature you are interested in. If you are switching from engine A to engine B and you knew how to do something in A, PEEL might be able to show you how to replicate that thing in B.


Figure 13: the same limited hinge joint in different engines

9) Keys summary

ESC	Quit
Return	Activate test (Menu Mode), End simulation (Simulation Mode)
C	Reset camera
+	Next camera pose
-	Previous camera pose
P	Enable/disable pause
O	Simulate one frame then pause
R	Enable/disable rendering
H	Show/hide help text
S	Save Excel file containing performance data
W	Enable/disable wireframe
X	Enable/disable wireframe overlay
Suppr	Delete picked object
Arrow keys	Test selection (Menu Mode), camera motion (Simulation Mode) or Vehicle control.

Appendix A: physics benchmarks for dummies

Benchmarking on PC is a black art. Benchmarking physics engines is even harder. Use the following notes to avoid the most basic mistakes.

Use the proper power options.

This is typically found in *Control Panel => System and security => Power Options*. Select the "High performance" power plan. Running benchmarks with the "Balanced" or "Power saver" plans produces unreliable results.
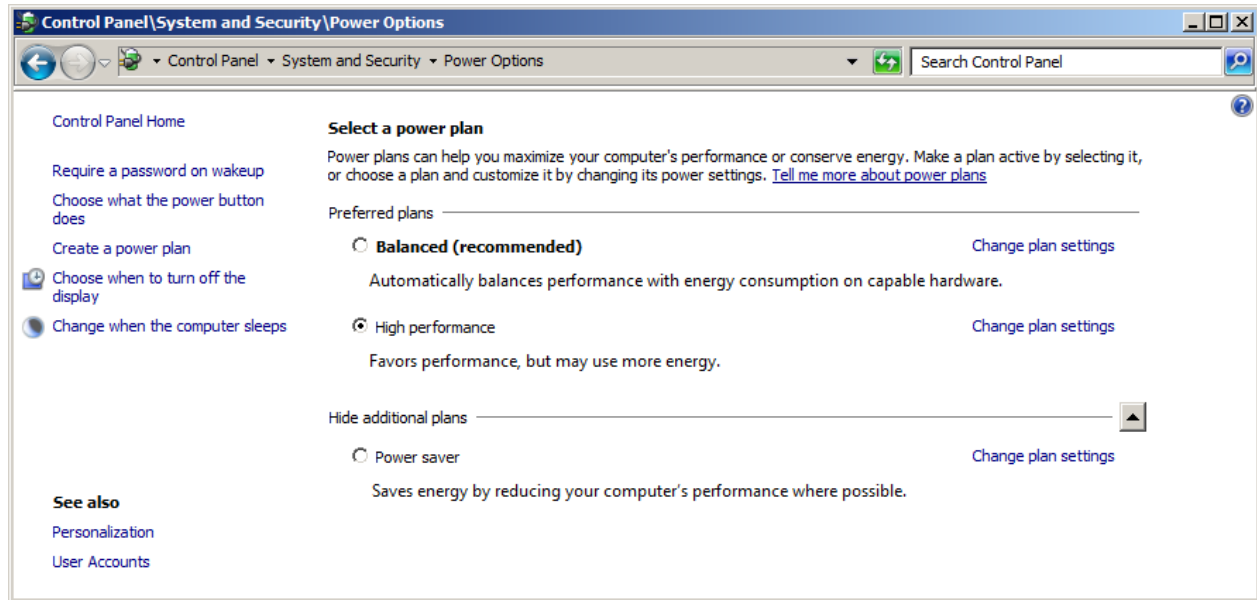

Figure 14: power options

Close all programs except PEEL. Unplug the internet.

Do not let programs like Outlook, Winamp, antivirus software, etc, run in the background. They can start random tasks at random times that will interfere with your benchmarks.

Ideally, start the Task Manager and kill all unnecessary processes. There are so many here that listing them all is impossible, but with some experience you should be able to know which ones can be killed, and which ones are worth killing.

It is of course very tedious to do this each time. So ideally you would take a radical step and use a dedicated PC with a fresh Windows installation and no internet connection. That is exactly what I do, and PEEL's benchmark results at home are a lot more stable than PEEL's benchmark results at work. Even when I do unplug the internet cable on the work PC…

For example Figure 15 shows the same scene captured at home and at work in "the same" conditions (boot the PC, run PEEL, do the benchmark). We see that the office PC is faster, but it also has a higher amount of background noise. And that amount can become a lot larger, as seen in Figure 12.
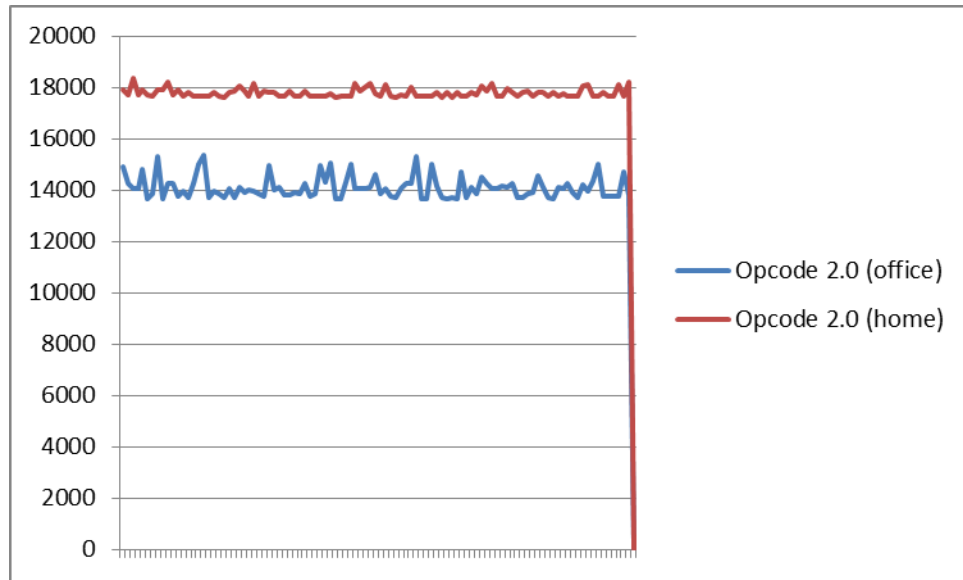
Figure 15: stable vs noisy benchmarks

Be aware of each engine's "empty" operating overhead.

In theory, when you run a physics update on an empty scene, all engines should take the same amount of time, i.e no time at all since there is nothing to do.

In practice, of course, this is not the case. PEEL's first test scene measures this operating cost. Figure 16 shows various results expressed in us.


Figure 16: operating overhead

Avoid benchmarks with just one object.

As a consequence, avoid running benchmarks with just a few objects or even a single object. The simulation time for just one object is likely to be lower than the engine's empty operating overhead, because the main internal algorithms are usually a lot more optimized than the glue code that connects them all together. Thus, such benchmarks actually measure this operating overhead more than anything else. While it is an interesting thing to measure, it does not reflect the engines' performance in real cases: the empty overhead is a constant time cost which is going to be lost in the noise of an actual game.

Thus, for example, it would be very wrong to run a benchmark with a single object and conclude that "engine A is faster than engine B" based on such results.

Try small scenes and large scenes.

Not all engines scale well. Some engines may be faster with small scenes, but collapse completely with large scenes – because large scenes have a tendency to expose O(N^2) parts of an engine.

Traditionally it is wise to "optimize for the worst case", so benchmarks involving large scenes tend to have a higher weight than those involving small scenes. Note that "small" and "large" are vague terms on purpose: a large scene in a game today might be considered a small scene in a game tomorrow. And at the end of the day, if it is fast enough for your game, it does not matter that an engine does not scale beyond that. It may matter for your next game though.

The point is: here again it is difficult to conclude from a limited set of benchmarks that "engine A is faster than engine B". You may have to refine your conclusions on a case-by-case basis.

Be aware of sleeping.

Virtually all physics engines have "sleeping" algorithms in place to disable work on non-moving, sleeping objects.

While the performance of an engine simulating sleeping objects is important, it is usually not the thing benchmarks should focus on. In the spirit of optimizing the worst case again, what matters more is the engine's performance when all these objects wake up: they must do so without killing the game's framerate.

Thus, PEEL typically disable sleeping algorithms entirely in its benchmarks, in order to capture the engines' 'real' performance figures. Unfortunately some physics engines may not let users disable these sleeping mechanisms, and benchmarks can appear biased as a result – giving an unfair advantage to the engines that put all objects to sleep.

Obviously, concluding that engine A (with sleeping objects) is faster than engine B (with non-sleeping objects) is foolish. Keep your eyes open for this in your experiments and benchmarks.

Be aware of solver iteration counts.

Most physics engines have a fast iterative solver that uses a default number of iterations. That default value may be different in each engine. For fair comparisons, make sure compared engines use the same number of iterations.

Alternatively, tweak the number of iterations in each engine until they all use roughly the same amount of time, then check which one produces the best simulation quality for the same CPU budget.

If a complex scene e.g. with joints does not work well by default in engine A, but works well with engine B, think about increasing the number of iterations for engine A. It might make it work while still remaining cheaper overall than engine B. And so on.

Comparing how engines behave out-of-the-box, with their default values, is only the tip of the iceberg.

Artificial benchmarks are not an actual game.

What works in the lab does not always work in the field. A good result in an artificial benchmark may not translate to a similarly good result in the final game. Good results in artificial benchmarks are just hints and good signs, not definitive conclusions. Take the results with the proverbial grain of salt.

Benchmarks are often artificial because they capture situations that would not actually happen in a game. At the same time, situations that would actually happen in a game often aren't complicated enough to expose significant differences between engine A and engine B, or they are too complicated to recreate in a benchmark environment.

Similarly, physics usually only takes a fraction of the game's frame. Thus, if engine A is "2X faster" than engine B in benchmarks, it does not mean that using engine A will make your game 2X faster overall. If your physics budget is 5% of the frame, even if you switch to an incredible physics engine that takes absolutely no time, you still only save 5% of the game's frame. Thus, it might actually be reasonable and acceptable to switch to a *slower* engine if it offers other benefits otherwise (better support, open source, etc).

Benchmarks are never "done".

There is always some possible scenario that you missed. There is always a case that you did not cover. There is maybe a different way to use the engine that you did not think about. There is always the possibility that an engine shining in all available benchmarks performs poorly in some other cases that were not captured.

There are more than 300 tests in PEEL, and still it only scratches the surface of what supported physics engines can do. Already though, in the limited set of available tests, no single engine always ends up "fastest". Sometimes engine A wins. Sometimes engine B wins.

Appendix B: list of supported physics engines

At time of writing, PEEL has plugins for:

- Bullet 2.79
- Bullet 2.81
- Bullet 2.82
- Bullet 2.85.1
- Newton 3.9
- Newton 3.12
- Newton 3.13
- Newton 3.14
- Havok 6.6.0
- Havok 2011_3_0
- Havok 2011_3_1
- Havok 2012_1_0
- Havok 2012_2_0
- Havok 2013_1_0
- Havok 2014_1_0
- ICE Physics
- NovodeX 2.1.1
- Opcode 1.3
- Opcode 2.0
- PhysX 2.8.4
- PhysX 3.1
- PhysX 3.2
- PhysX 3.3.0
- PhysX 3.3.1
- PhysX 3.3.2
- PhysX 3.3.4
- PhysX 3.4 (*many* versions)
- GRB (GPU rigid bodies)

Writing and maintaining the PEEL plugins for each of these engines takes a lot of time and effort. It goes without saying that it also requires a good knowledge of each of these engines' API, which is something the author does not necessarily have. Thus, these implementations are a "best effort", and it is possible that some of them are suboptimal or incomplete. A list of current known issues is maintained in a separate text file in the PEEL distribution. The author strongly encourages people familiar with the above APIs to look up the code, double-check it, *and submit improvements or bugfixes*.

Note that PEEL can also be used to evaluate collision libraries like Opcode, which are not "physics engines" per se, but they offer similar collision detection features.

Credits & thanks

PEEL and most of the PINT plugins were created by Pierre Terdiman.

The Newton plugin was written by Julio Jerez.

Thanks to Dirk Gregorius for his comments and feedback.